

ML@LSE - Bootcamp 5: Feedforward Neural Networks

Introduction: What are neural networks used for?

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. **Should I go to Brighton this weekend?**
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. FNN is a network of perceptrons

B- Fitting the network

1. Minimizing the Cost Function

- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

We've studied methods of supervised and unsupervised learning. Neural networks can perform both of these tasks. It also belongs to another class of ML: deep learning. Within supervised learning, it can be used for classification and regression.

Neural networks are very polyvalent!

They are used in:

- Image recognition
- Speech recognition
- Credit risk modeling
- Time-series forecasting
- Algorithmic trading
- etc...

The intuition of NN

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. **Framing this with a perceptron**

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. FNN is a network of perceptrons

B- Fitting the network

1. Minimizing the Cost Function

- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

Should I go to Brighton this weekend?

➤ How do you make such a decision? What parameters do you consider?

- 1. How many friends are willing to go? → x_1
- 2. Will it be sunny? → x_2
- 3. How much work have I got to do? → x_3
- 4. How much money do I have → x_4

How to make the decision? Weight each parameter and decide!

Go if: $w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 \geq b$

$$w \cdot x \geq b$$

$$y(w \cdot x - b) = \begin{cases} 1 & \text{if } w \cdot x - b \geq 0 \\ 0 & \text{if } w \cdot x - b \leq 0 \end{cases} \quad \text{This is called the activation function}$$

The intuition of NN

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. **Framing this with a perceptron**

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. FNN is a network of perceptrons

B- Fitting the network


1. Minimizing the Cost Function

- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

We can represent this schematically:




The sigmoid activation function

The activation function $y(w \cdot x - b)$ is binary (returns either 0 or 1).

In the following we'll use a more subtle activation function that returns any number between 0 and 1. You can think of these numbers as probabilities (if the activation function returns a number close to one, the probability that I will have a good time in Brighton is high).

This activation function is the sigmoid function: $\sigma(w \cdot x - b) = \frac{1}{1 + e^{-(w \cdot x - b)}}$



A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. **The sigmoid function**
- b. FNN is a network of perceptrons

B- Fitting the network

1. Minimizing the Cost Function


- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

FNN as a network of percetrons

A feedforward neural network can be seen as a network of percetrons:



A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. **FNN is a network of perceptrons**

B- Fitting the network

1. Minimizing the Cost Function

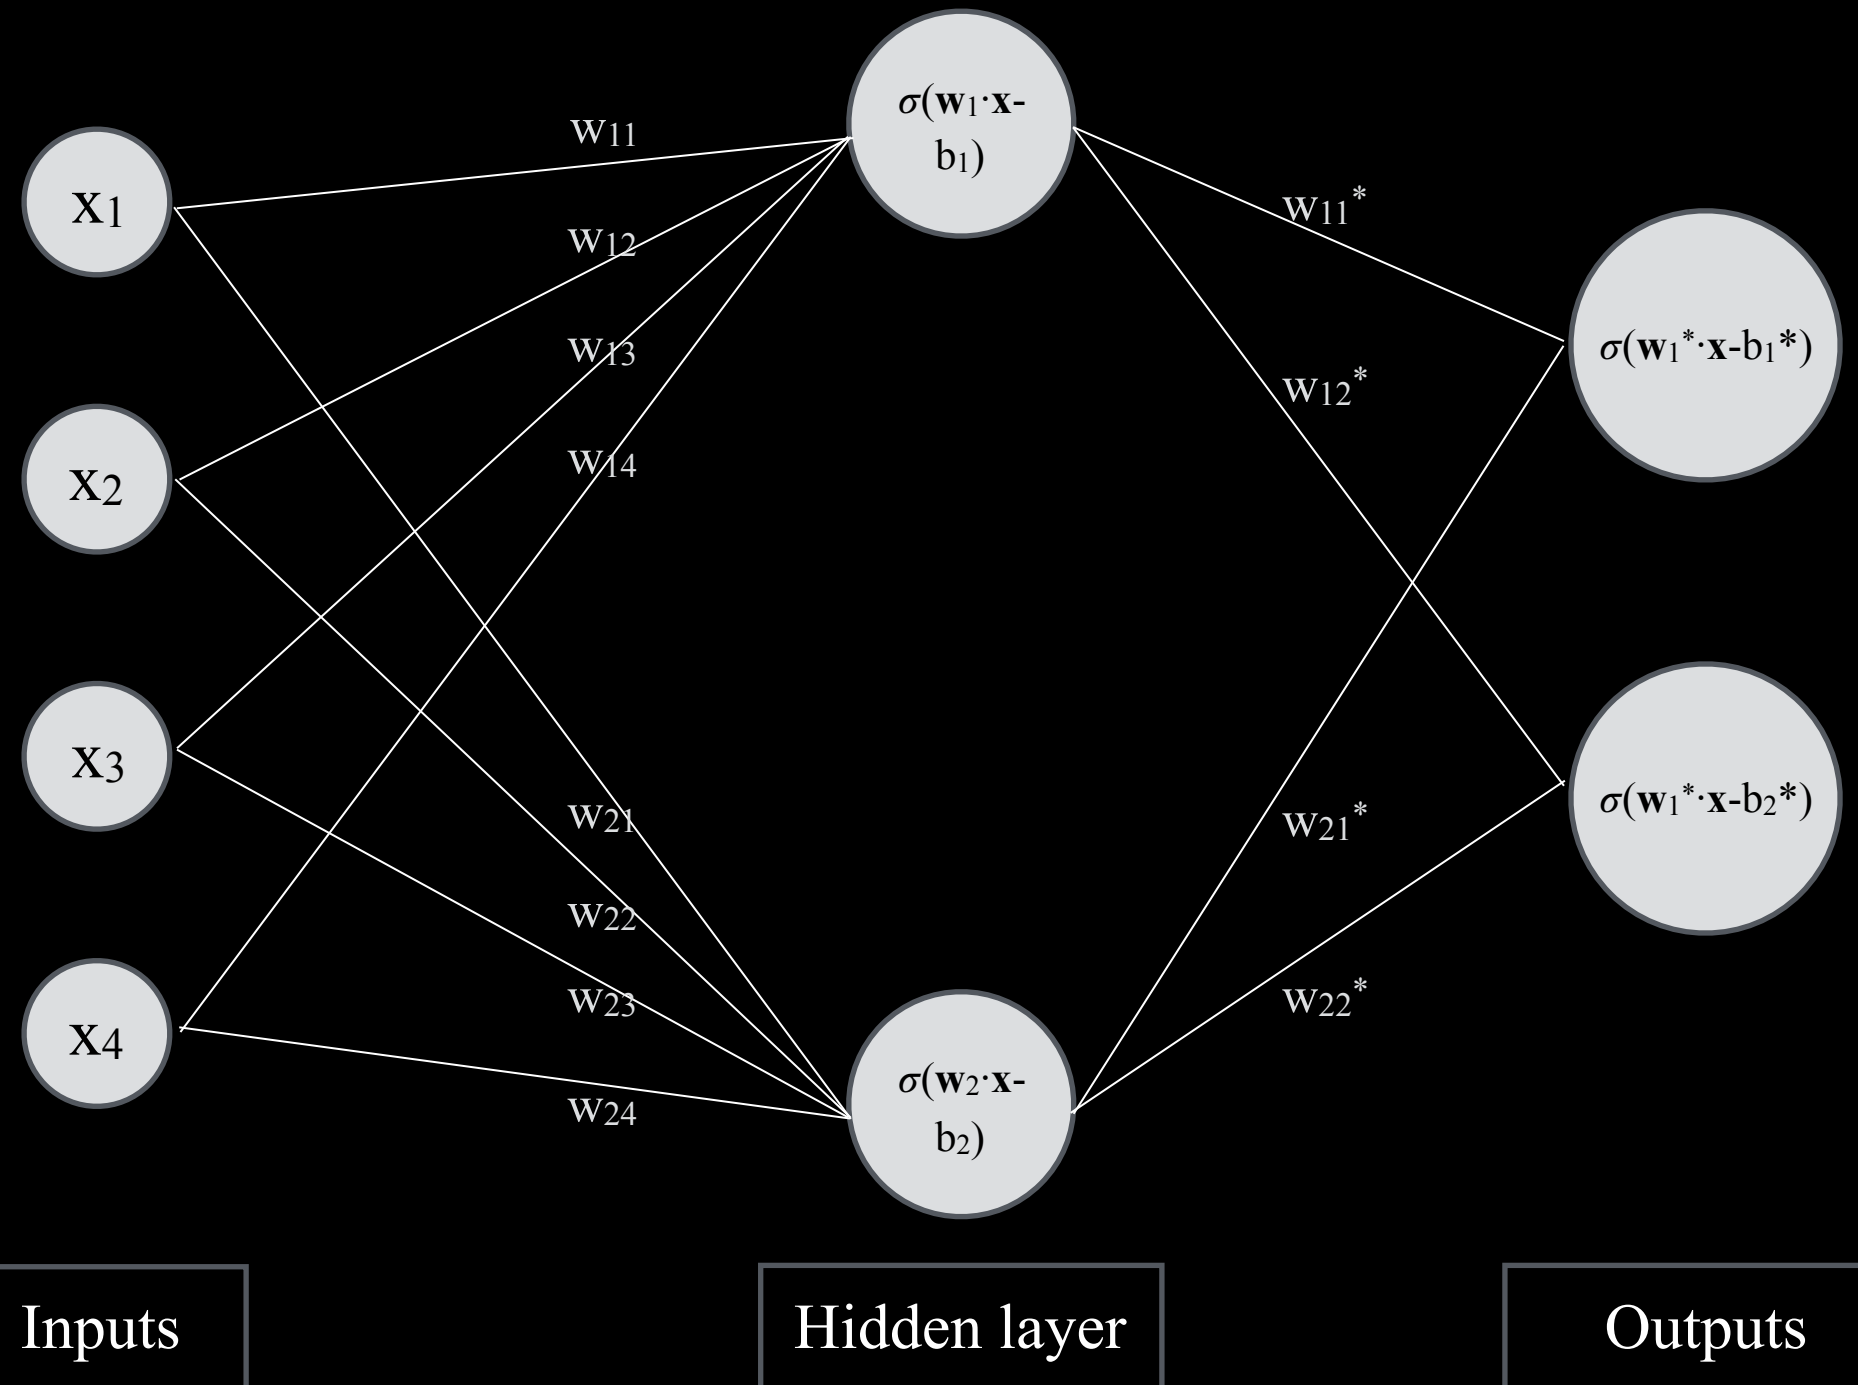
- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

FNN as a network of perceptrons

A simple neural network with one hidden layer:



A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- Should I go to Brighton this weekend?
- Framing this with a perceptron

2. Feedforward Neural Network structure

- The sigmoid function
- FNN is a network of perceptrons**

B- Fitting the network

1. Minimizing the Cost Function

- The Cost function of a network
- Gradient descent

2. Fitting the network using backpropagation

- Closer look at the gradient
- Backpropagation

FNN as a network of perceptrons

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. FNN is a network of perceptrons**

B- Fitting the network

1. Minimizing the Cost Function

- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

Hence a network is a sort of function $f(\mathbf{x})$ with inputs and outputs.
(The true mathematical term is a correspondence, as the network can have several outputs)

Importantly, this $f(\mathbf{x})$ has many parameters, which are?

➤ The weights and biases are the parameters of the network!

As we've seen many times now, machine learning oftentimes consists in shaping a function properly, using the training data, so that this function has some predictive power.

With a neural network, we will “shape the function” by adjusting the weights and biases properly.

A practical example

Let's look at a practical example:



➤ How can we use our neural network to classify these shapes?

What is the output?

- Probabilities assigned to each digit 0, 1, ..., 9

What is the input?

- These numbers are represented by a grid of 28 by 28 pixels, whose brightness varies from 0 to 1. Hence the input will be the brightness of each of the 784 pixels.

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. **FNN is a network of perceptrons**

B- Fitting the network

1. Minimizing the Cost Function

- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

A practical example

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- Should I go to Brighton this weekend?
- Framing this with a perceptron

2. Feedforward Neural Network structure

- The sigmoid function
- FNN is a network of perceptrons**


B- Fitting the network

1. Minimizing the Cost Function

- The Cost function of a network
- Gradient descent

2. Fitting the network using backpropagation

- Closer look at the gradient
- Backpropagation



A practical example

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. **FNN is a network of perceptrons**

B- Fitting the network

1. Minimizing the Cost Function


- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation



784



Cost

How do we fit the parameters?

➤ As usual, using a cost function:

Predicted	Actual
0.011324	0
0.000023	0
0.000453	0
0.000006	0
0.004300	0
0.003032	0
0.000455	0
0.000998	0
0.000067	0
0.988976	1

Take the sum of square differences:

$$(0.011324 - 0)^2 + \dots + (0.988976 - 1)^2$$

$$= \sum_{n=1}^N (\hat{y}_n - y_n)^2$$

Suppose we have sample of S written digits

$$Cost = \frac{1}{S} \sum_{s=1}^S \left(\sum_{n=1}^N (\hat{y}_{sn} - y_{sn}) \right)$$

$$Cost = C(\mathbf{w}, \mathbf{b})$$

We want to solve:

$$\min_{\mathbf{w}, \mathbf{b}} C(\mathbf{w}, \mathbf{b})$$

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. FNN is a network of perceptrons

B- Fitting the network

1. Minimizing the Cost Function

- a. **The Cost function of a network**
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

Cost is a complicated function!


How to solve $\min_{\mathbf{w}, \mathbf{b}} C(\mathbf{w}, \mathbf{b})$??

$$\frac{\partial C(\mathbf{w}, \mathbf{b})}{\partial \mathbf{w}} = 0, \frac{\partial C(\mathbf{w}, \mathbf{b})}{\partial \mathbf{b}} = 0$$

Maybe not, C has too many parameters!

What solution?

Let's use an analogy: imagine you're in Scotland on the top of a hill. You want to reach a valley but there's fog and you can't see further than 10 meters away. What is your best strategy?



A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. FNN is a network of perceptrons

B- Fitting the network

1. Minimizing the Cost Function

- a. The Cost function of a network
- b. **Gradient descent**

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation



Gradient descent

Gradient descent:

Let \mathbf{W} represent the set of both weights \mathbf{w} and biases \mathbf{b} , i.e. \mathbf{W} is the set of all parameters in the network that we want to optimize.

- Start with a random \mathbf{W}_0 , and then update it (go downhill) according to:

$$\mathbf{W}_{n+1} = \mathbf{W}_n - \eta \nabla C(\mathbf{W}_n), \quad \eta > 0$$

This equation means: “wherever you are now, go down as steep as you can.”

An approximation of this method, stochastic gradient descent (SGD), is usually use.

3D model

$$\nabla C(\mathbf{W}) = \begin{bmatrix} \frac{\partial C}{\partial w_1} \\ \frac{\partial C}{\partial w_2} \\ \dots \\ \frac{\partial C}{\partial w_m} \\ \frac{\partial C}{\partial b_1} \\ \dots \\ \frac{\partial C}{\partial b_k} \end{bmatrix}$$

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- Should I go to Brighton this weekend?
- Framing this with a perceptron

2. Feedforward Neural Network structure

- The sigmoid function
- FNN is a network of perceptrons

B- Fitting the network

1. Minimizing the Cost Function

- The Cost function of a network
- Gradient descent**

2. Fitting the network using backpropagation

- Closer look at the gradient
- Backpropagation

Partial derivatives

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. FNN is a network of perceptrons

B- Fitting the network


1. Minimizing the Cost Function

- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

$$\nabla C(\mathbf{W}) = \begin{bmatrix} \frac{\partial C}{\partial w_1} \\ \frac{\partial C}{\partial w_2} \\ \dots \\ \frac{\partial C}{\partial w_m} \\ \frac{\partial C}{\partial b_1} \\ \dots \\ \frac{\partial C}{\partial b_k} \end{bmatrix}$$



Backpropagation

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- Should I go to Brighton this weekend?
- Framing this with a perceptron

2. Feedforward Neural Network structure

- The sigmoid function
- FNN is a network of perceptrons

B- Fitting the network

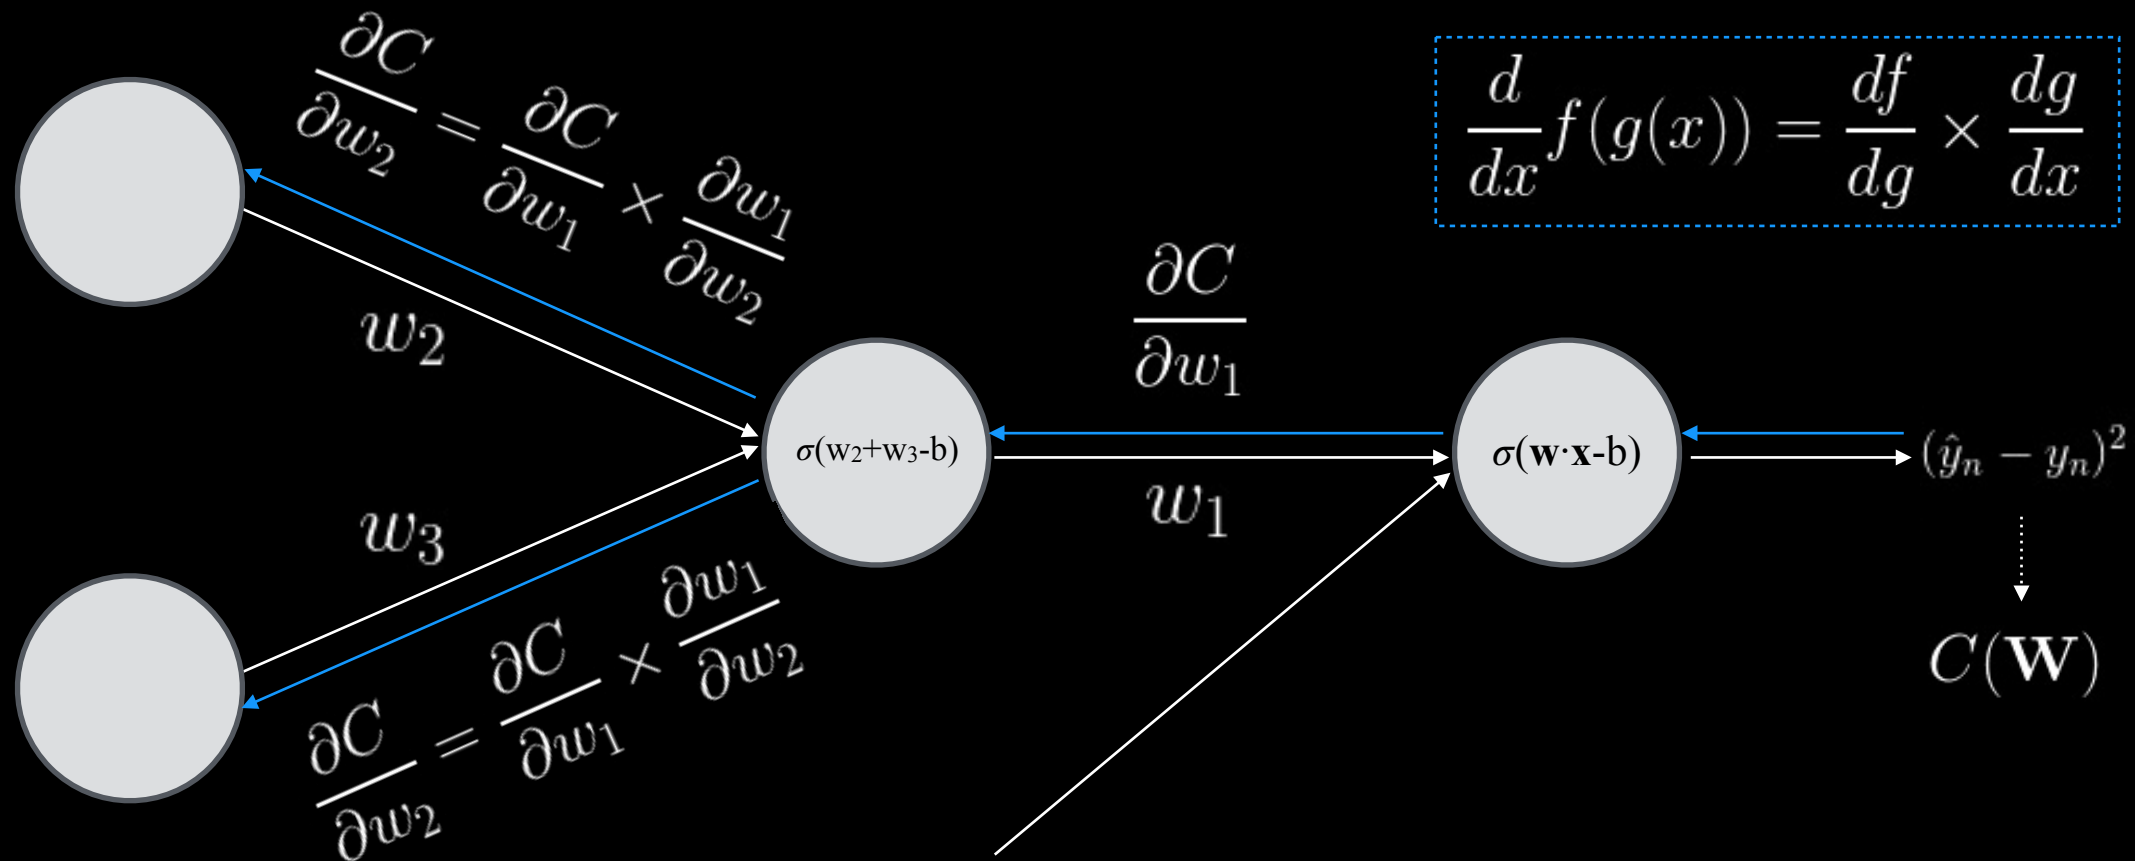

1. Minimizing the Cost Function

- The Cost function of a network
- Gradient descent

2. Fitting the network using backpropagation

- Closer look at the gradient
- Backpropagation**

$$\nabla C(\mathbf{W}) = \begin{bmatrix} \frac{\partial C}{\partial w_1} \\ \frac{\partial C}{\partial w_2} \\ \dots \\ \frac{\partial C}{\partial w_m} \\ \frac{\partial C}{\partial b_1} \\ \dots \\ \frac{\partial C}{\partial b_k} \end{bmatrix}$$



Backpropagation

A- The structure of Neural Networks

1. An intuitive example of a Neural Network

- a. Should I go to Brighton this weekend?
- b. Framing this with a perceptron

2. Feedforward Neural Network structure

- a. The sigmoid function
- b. FNN is a network of perceptrons

B- Fitting the network

1. Minimizing the Cost Function

- a. The Cost function of a network
- b. Gradient descent

2. Fitting the network using backpropagation

- a. Closer look at the gradient
- b. Backpropagation

<http://www.emergentmind.com/neural-network>

Thanks for coming!